

A APPENDIX

A.1 Proof of Theorem 3.1

We begin by introducing some notation from prior work [1]. First, recall that a degree sequence for a relation R and variable V is a vector $f_{R,V}$ such that $f_{R,V,i}$ is the frequency of the i th most frequent value, and $F_{R,V,i} = \sum_{j=1,i} f_{R,V,j}$ is the cumulative degree sequence (CDS). We represent relations R as tensors M_R with one dimension for each variable present in the relation. The value at a particular entry $M_{i_1, \dots, i_{|V_R|}}$ is equal to the frequency of that tuple, $(i_1, \dots, i_{|V_R|})$, in the relation R . Note that entries can be zero if that tuple does not appear in R . Additionally, the discrete derivative and integral of a tensor M on a variable V is defined as,

Definition A.1.

$$(\Delta_V M)_{v_i} = M_{v_i} - M_{v_i-1} \quad (1)$$

$$(\Sigma_V M)_{v_i} = \sum_{j=1, v_i} M_j \quad (2)$$

A tensor M is *consistent* with a set of degree sequences f_R if the following is true,

$$\left(\sum_{V' \neq V} M \right)_i \leq f_{R,V,i} \quad \forall i \in \mathbb{D}_V, V \in V_R$$

Briefly, this means that if we contract the tensor down to a single dimension, then the resulting vector is less than the DF $F_{R,V}$ at all points. The set \mathcal{M}_{f_R} is the set of tensors consistent with f_R .

Lastly, we define the *value tensor*, E^{F_R} , and use it to explicitly define the *worst-case tensor*, C^{F_R} .

Definition A.2. The *value tensor*, $E^{F_R} \in \mathbb{R}_+^{[n]}$, is defined by the following linear optimization problem:

$$\begin{aligned} \forall \mathbf{m} \in [n] : \quad E_{\mathbf{m}}^{F_R} &\stackrel{\text{def}}{=} \text{Maximize: } \sum_{s \leq \mathbf{m}} M_s \\ \text{Where: } \mathbf{M} &\in \mathcal{M}_{f_R} \end{aligned} \quad (3)$$

The *worst-case tensor*, $C^{F_R} \in \mathbb{R}^{[n]}$, is defined as:

$$C^{F_R} \stackrel{\text{def}}{=} \Delta_{V_1} \cdots \Delta_{V_d} E^{F_R} \quad (4)$$

Or, equivalently,

$$\Sigma_{V_1} \cdots \Sigma_{V_d} C^{F_R} = E^{F_R} \quad (5)$$

Note that this worst-case tensor is equivalent to the worst-case instance, $W(R)$ of a relation R , as depicted in Figure 2.

Getting back to Theorem 3.1 of this work, we start by proving that it holds for star queries before expanding to all Berge-acyclic queries. Consider the following query,

$$Q_{STAR} = R(V_1, \dots, V_d) S_1(V_1) \cdots S_d(V_d)$$

If \mathbf{M} is the count tensor of the relation R and $\mathbf{a}^{(V_i)}$ is the count tensor of S_i , in this case a simple non-increasing vector, then we can express the query size as follows,

$$|Q_{STAR}(D)| = \mathbf{M} \cdot \mathbf{a}^{(V_1)} \cdots \mathbf{a}^{(V_d)}$$

Given this notation, we consider part of Theorem 3.2 from [1].

Theorem A.3. [Thm. 3.2 from [1]] Let f_R be the set of degree sequences as above, and let V, C defined by (3)-(5). Then:

(1) We can define the *value tensor* as follows,

$$\forall \mathbf{m} \in [n] : \quad E_{\mathbf{m}}^{F_R} = \min (F_{R,V_1}(m_1), \dots, F_{R,V_d}(m_d)) \quad (6)$$

(2) For any non-increasing vectors $\mathbf{a}^{(V_p)} \in \mathbb{R}_+^{[n_p]}$, $p = 2, d$, the vector $C^{F_R} \cdot \mathbf{a}^{(V_2)} \cdots \mathbf{a}^{(V_d)}$ is in $\mathbb{R}_+^{[n_1]}$ and non-increasing.

(3) For all count tensors \mathbf{M}_R , and all non-increasing vectors $\mathbf{a}^{(X_1)} \in \mathbb{R}_+^{[n_1]}, \dots, \mathbf{a}^{(X_d)} \in \mathbb{R}_+^{[n_d]}$:

$$\mathbf{M}_R \cdot \mathbf{a}^{(V_1)} \cdots \mathbf{a}^{(V_d)} \leq C^{F_R} \cdot \mathbf{a}^{(V_1)} \cdots \mathbf{a}^{(V_d)} \quad (7)$$

Directly implying,

$$|Q_{STAR}(D)| \leq |Q_{STAR}(W(s(D)))| \quad (8)$$

Let \hat{F}_R be an upper bound of F_R , i.e. $\hat{F}_{R,V}(i) \geq F_{R,V}(i) \forall V \in V_R, i$, and define $\hat{f}_{R,V} = \Delta_V \hat{F}_{R,V}$ and \hat{f}_R as the set of these degree sequences which, as specified in Theorem 3.1, must be non-increasing. Further, note that item 1 and item 1 relies only on the properties of the worst-case instance's inherent structure, so it immediately applies to $C^{\hat{F}_R}$.

Based on the above, we can prove the following lemma,

Lemma A.4. For all non-increasing vectors $\mathbf{a}^{(V_1)} \in \mathbb{R}_+^{[n_1]}, \dots, \mathbf{a}^{(V_d)} \in \mathbb{R}_+^{[n_d]}$:

$$C^{f_R} \cdot \mathbf{a}^{(V_1)} \dots \mathbf{a}^{(V_d)} \leq C^{\hat{F}_R} \cdot \mathbf{a}^{(V_1)} \dots \mathbf{a}^{(V_d)} \quad (9)$$

Directly implying,

$$|Q_{STAR}(W(s(D)))| \leq |Q_{STAR}(W(\Delta \hat{S}))| \quad (10)$$

PROOF. Following the original proof of item 2, we begin by simplifying the problem using 1-0 vectors. In particular, let $\mathbf{b}^{(m)} \in \mathcal{R}^n$ be the vector with m 1's followed by $n - m$ 0's. Because the \mathbf{a}^{V_i} are non-increasing integral vectors, they can be represented as a sum of 1-0 vectors, so it suffices to consider the case where each of them is a 1-0 vector. In this case, the problem description becomes,

$$C^{f_R} \cdot \mathbf{b}^{(m_1)} \dots \mathbf{b}^{(m_d)} \leq C^{\hat{F}_R} \cdot \mathbf{b}^{(m_1)} \dots \mathbf{b}^{(m_d)}$$

Multiplying against $\mathbf{b}^{(m)}$ is the same as summing over the first m indices, so this can be alternatively expressed as,

$$\Sigma_{m_1} \dots \Sigma_{m_d} C^{f_R} \leq \Sigma_{m_1} \dots \Sigma_{m_d} C^{\hat{F}_R}$$

Considering the definition of the value tensor $E_m^{f_R}$, we can rephrase this as follows where $\mathbf{m} = (m_1, \dots, m_d)$,

$$E_m^{f_R} \leq E_m^{\hat{F}_R}$$

Lastly, we insert the alternative definition of $E_m^{f_R}$ provided in item 1 and the fact that each \hat{F}_{R,V_i} is an upper bound of F_{R,V_i} to prove the lemma,

$$\min(F_{R,V_1}(m_1), \dots, F_{R,V_d}(m_d)) \leq \min(\hat{F}_{R,V_1}(m_1), \dots, \hat{F}_{R,V_d}(m_d))$$

□

To prove that this can be extended to general queries, we rely on more theory from [1].

Theorem A.5. Implied by Thm. 4.2 of [1] If the following is true for a set of database instances, \mathcal{D} ,

- (1) For any non-increasing vectors $\mathbf{a}^{(V_p)} \in \mathbb{R}_+^{[n_p]}, p = 2, d$, the vector $C^{f_R} \cdot \mathbf{a}^{(V_2)} \dots \mathbf{a}^{(V_d)}$ is in $\mathbb{R}_+^{[n_1]}$ and non-increasing.
- (2) For all relations $R \in \mathcal{D}$ with count tensor M_R , and all non-increasing vectors $\mathbf{a}^{(X_1)} \in \mathbb{R}_+^{[n_1]}, \dots, \mathbf{a}^{(X_d)} \in \mathbb{R}_+^{[n_d]}$:

$$M_R \cdot \mathbf{a}^{(V_1)} \dots \mathbf{a}^{(V_d)} \leq C^{f_R} \cdot \mathbf{a}^{(V_1)} \dots \mathbf{a}^{(V_d)} \quad (11)$$

Then, for any Berge-acyclic query, Q ,


$$|Q(D)| \leq |Q(W(s))| \quad \forall D \in \mathcal{D} \quad (12)$$

The first immediately holds for $C^{\hat{F}_R}$ from Theorem A.3 while the latter holds for the set of all database instances, \mathcal{D} , such that $D \models S$ due to Lemma A.4.


B DISAGGREGATED RESULTS

In this section, we provide alternate versions of Figures 6, 7, and 9.a which are disaggregated by benchmark.


C RELATIVE ERROR BY JOIN SIZE




(a) Top 40 Query Runtimes - JOBLight



(b) Top 40 Query Runtimes - JOBLightRanges



(c) Top 40 Query Runtimes - JOBM



(d) Top 40 Query Runtimes - Stats

Figure 1: Top Query Runtime Graphs

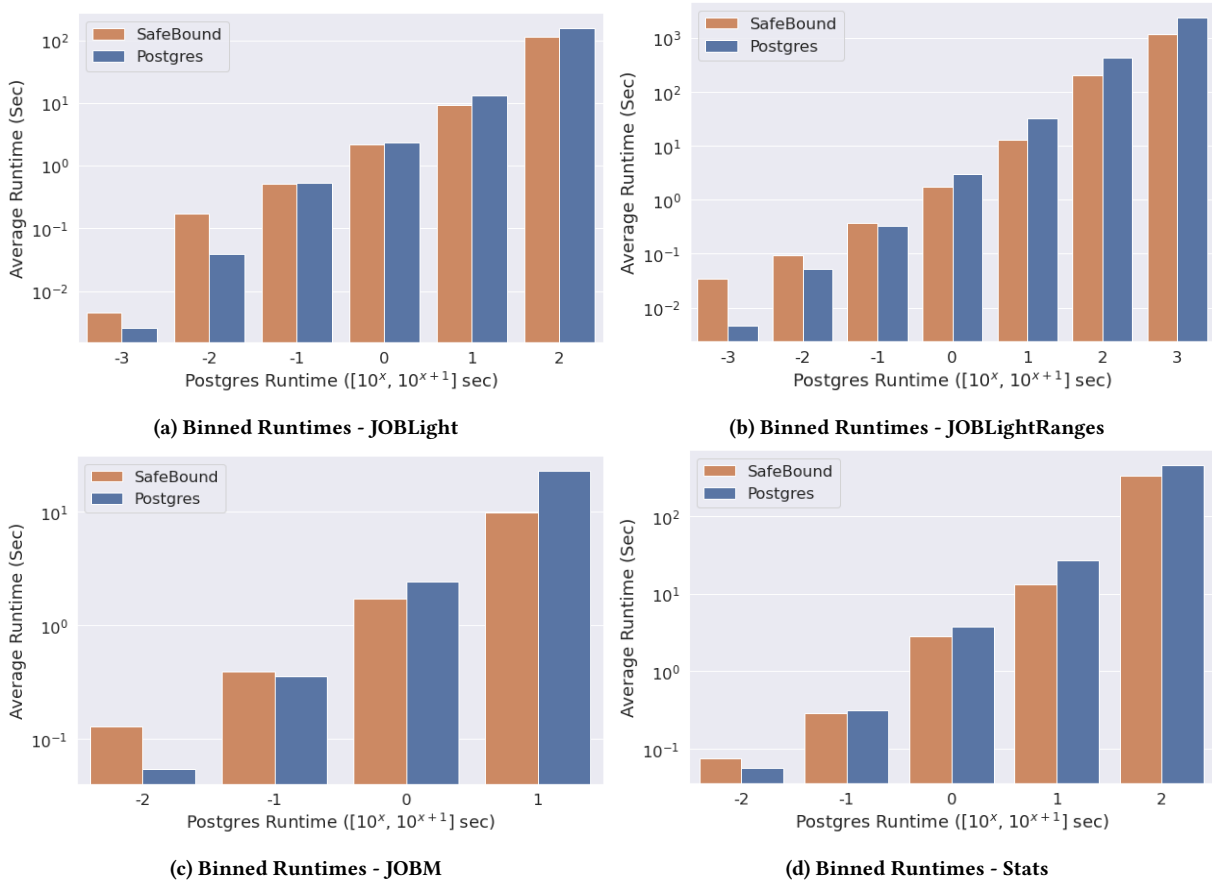


Figure 2: Binned Runtime Graphs

REFERENCES

- [1] Kyle Deeds, Dan Suciu, Magda Balazinska, and Walter Cai. 2022. Degree Sequence Bound For Join Cardinality Estimation. *CoRR* abs/2201.04166 (2022). [arXiv:2201.04166](https://arxiv.org/abs/2201.04166) <https://arxiv.org/abs/2201.04166>

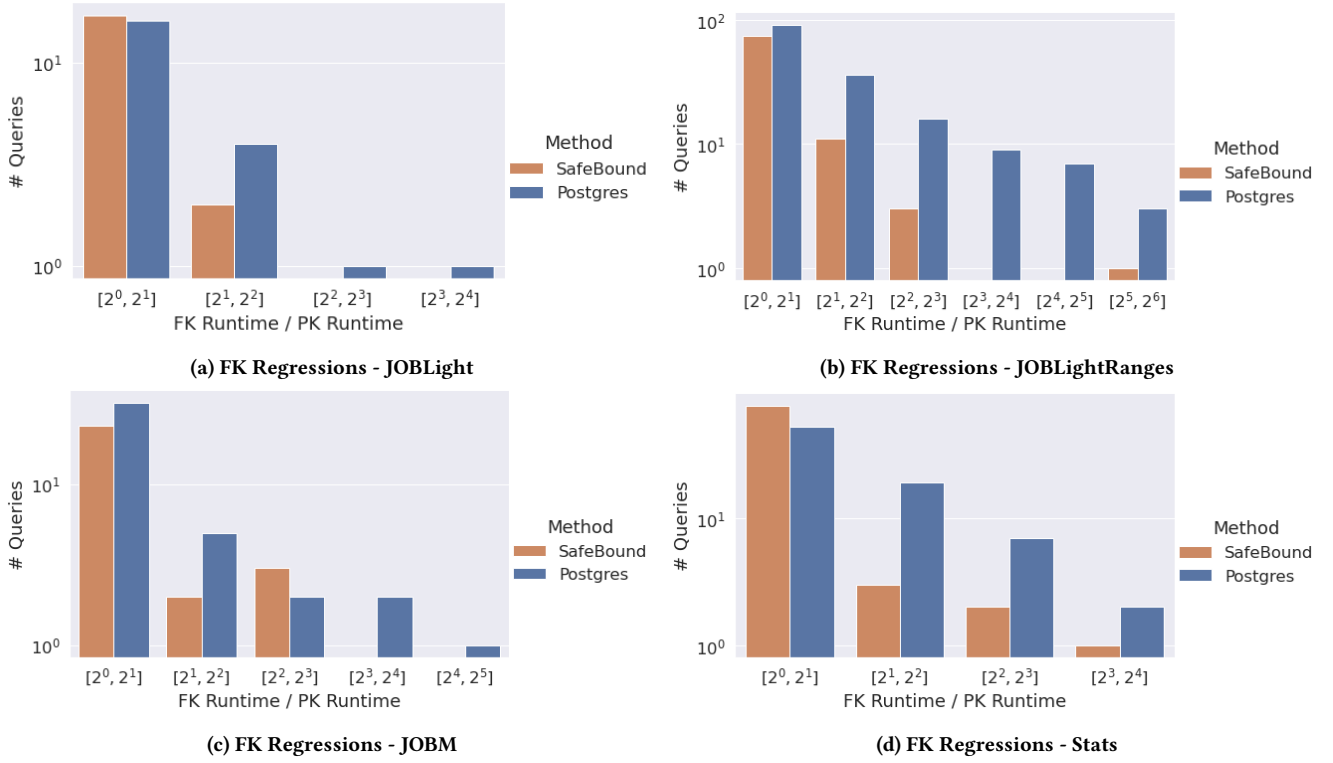




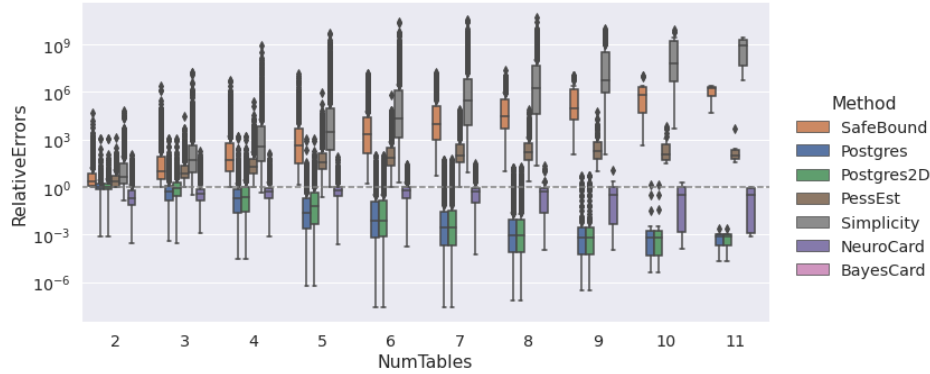
Figure 3: FK Regressions Graphs



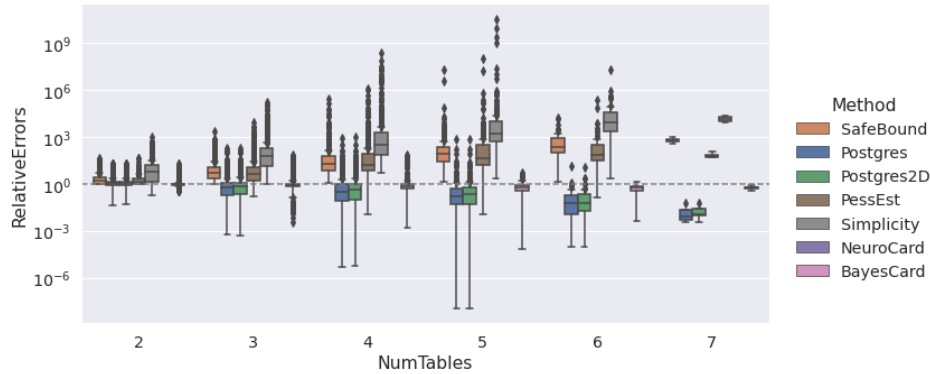
(a) Relative Error By Join Size - JOBLight



(b) Relative Error By Join Size - JOBLightRanges



(c) Relative Error By Join Size - JOBM



(d) Relative Error By Join Size - Stats

Figure 4: Relative Error By Join Size Graphs